

Expert Parallelism, from scratch

A custom `expert_parallel` module I built on top of `picotron`

From a naive all-to-all to DeepEP, tiled pipelines, FP8, LatentMoE — and what it takes to actually *train* at scale


extending picotron (DP·PP·CP·TP) with a 5th axis — expert parallelism — and a stack of MoE comm optimizations

deep dive: autograd · CUDA streams · Triton kernels · the bandwidth lesson

What we are optimizing

A **Mixture-of-Experts** FFN: a router picks the top- k of E experts per token.

$$y_t = \sum_{e \in \text{top-}k(t)} g_{t,e} \text{FFN}_e(x_t) + \text{shared}(x_t), \quad g_{t,e} = \text{softmax}(W_g x_t)_e$$




Only k of E experts run per token (sparse). The router is trained; the shared expert is

EP: tokens travel to their expert and back

Expert Parallelism: shard the E experts across ranks — rank r owns $[r \cdot L, (r+1) \cdot L]$,
 $L = E/\text{ep_size}$.

A token routed to a remote expert needs **two all-to-alls**: dispatch out, combine back.



EP: the 5th axis I added to picotron

Stock picotron is `DP × PP × CP × TP` (4D). I extended its process-group manager to a 5D grid `DP × PP × CP × EP × TP` and added the MoE layer that uses it.

Tensor	Sharded over	Grad sync
Expert weights <code>FFN_e</code>	EP	none (each rank owns distinct experts)
Router <code>W_g</code> , shared expert, norms	replicated	all-reduce over <code>cp_dp</code>

Design constraints I kept:

- **Readable** `torch.distributed` collectives (picotron spirit)
- **Bit-exact** across every optimization path (tested vs the naive baseline, `grad_diff ≤ 5e-7`)
- Every knob is **optional** and safe to leave on (auto-fallback)


The challenge: sparsity breaks the dense playbook

Dense: every param fires every token → params and compute scale together; sharding keeps comm small.

MoE: only k of E experts fire → total params grow far faster than per-token compute = a **parameter–compute mismatch**.

“ DeepSeek-V3: **685B** total / **37B** active = **18x** · Kimi-K2 = **31x** ”

That asymmetry creates **three tightly-coupled walls** (NVIDIA *Megatron-Core MoE*, [arXiv:2603.07685](https://arxiv.org/abs/2603.07685)):



The Three Walls — what actually hurts

Memory

All **E** experts' **params + grads + optimizer** stay resident, though only **k** fire. Relieving it just **moves the cost** (shard→BW, recompute→FLOPs, offload→PCIe). Dynamic routing → **memory spikes**.

Communication

dispatch+combine all-to-all,
 $\approx T \cdot k \cdot h \cdot (EP-1) / EP$ each way. As EP grows it leaves NVLink for **inter-node links (~10× slower)**, and sparse compute leaves **little to overlap → 20–60%** of the step.

Compute efficiency

- **small GEMMs** from fine-grained experts → idle SMs
- **routing/permute** ~9% of layer time
- **load imbalance** → some experts idle
- **host overhead**: many kernel launches; droptess dynamic shapes force host-device syncs

How this talk knocks down the walls

Wall	Levers in this module
Communication	overlap (shared-expert · tiled · DeepEP) + volume (FP8 · LatentMoE)
Memory	static-memory capacity dispatch (fixed buffers)
Compute	grouped-GEMM megakernel · load balancing (aux loss + loss-free bias)

We start with the **Communication Wall** — the one EP creates — then return to **Memory** and

Compute in *Beyond speed*. Throughline: **measure** $\rho = t_{\text{comm}}/t_{\text{compute}}$, then spend effort where the wall actually is.

Knocking down the Communication Wall: two levers

1. Hide the comm (overlap)

Run useful compute *while* the all-to-all is in flight.

- shared-expert overlap (side stream)
- tiled pipeline (fwd **and** bwd)
- DeepEP kernels (near-zero SM)

→ helps **only** when comm is a real fraction of the step.

2. Send fewer bytes (volume)

Shrink the payload itself.

- FP8 (E4M3) dispatch
- LatentMoE (route/compute in a small latent dim)

→ helps even with **nothing left to hide behind**.

Orthogonal. They stack: `latent + fp8` cuts bytes 8×, and overlap hides whatever remains.

Step 0 – the naive baseline

dispatch → experts → combine, blocking, fully differentiable.

```
def _moe_naive(self, tokens, topk_idx, topk_weights):
    routed, expert_idx, w = self._expand_routed(tokens, topk_idx, topk_weights)
    order, inv, in_list, out_list, local_expert = self._dest_plan(expert_idx)
    recv      = all_to_all(routed[order], out_list, in_list, group)  # dispatch
    recv_out  = self._run_local_experts(recv, local_expert)          # grouped GEMM
    combined  = all_to_all(recv_out, in_list, out_list, group)        # combine
    y = self._combine_topk(combined[inv], w, num_tokens)              # weight + sum top-k
    return self._finalize(y, shared)
```


The all-to-all itself is one autograd `Function` – **backward of an a2a is the reverse a2a:**

```
class _AllToAll(torch.autograd.Function):
    def forward(ctx, group, x, out_splits, in_splits): ...  # all_to_all_single(out, x, out, in)
    def backward(ctx, g): ...                               # all_to_all_single(gin, g, in, out) ← swapped
```

The question that decides everything

How much of the step is communication? $\rho = t_{\text{comm}} / t_{\text{compute}}$

Naive: everything serial on the critical path



As EP scales, the two a2a's go off-node and grow:

$\rho = t_{\text{comm}} / t_{\text{compute}}$ grows as the two orange bars dominate the step.

- $\rho \gg 1$ (cross-node IB/RoCE, large EP) → **overlap is gold** · $\rho \ll 1$ (NVLink island) → **nothing to hide**

We *measure* ρ on three fabrics and watch the same code flip from **1.6× win** to **no-op**.

Opt 1 – shared-expert overlap (the orthogonal one)

DeepSeek-style models add an **always-on shared expert** = pure local compute, **no comm**.

Idea: run it on a **side CUDA stream** so it overlaps the routed path's all-to-all.


```
def _shared_start(self, tokens):
    if not (self.ep_overlap and self.ep_world_size > 1 and tokens.is_cuda):
        return self.shared_expert(tokens), None          # serial (CPU/gloo, 1 rank)
    self._shared_stream.wait_stream(torch.cuda.current_stream()) # tokens ready
    with torch.cuda.stream(self._shared_stream):
        out = self.shared_expert(tokens)                  # runs concurrently with dispatch/combine
    return out, self._shared_stream

def _finalize(self, y, shared):
    out, stream = shared
    if stream is not None:
        torch.cuda.current_stream().wait_stream(stream) # join before the add
    return y + out
```

Key design choice: this is *not* a 4th path. It composes with **naive / tiled / deeppep** alike.

Opt 1 – deep notes (streams + autograd)

Shared-expert overlap: hide comm on a side CUDA stream




- **Correctness:** `wait_stream` on entry + before the add; reads of `tokens` on two streams don't conflict.
- **Backward:** autograd replays the shared-expert backward on the side stream (one benign `AccumulateGrad` stream-mismatch warning) – verified **bit-exact** (`out_diff = grad_diff = 0` on 2× A100 NCCL).
- **Best on DeepEP:** its comm uses ~0 SMs, so the GPU is free for the shared FFN.

`num_shared_experts = N` merges into one MLP of width `N · intermediate` (= sum of N experts, one GEMM).

Opt 2 – tiled pipeline (forward)

Split routed tokens into N tiles; software-pipeline the three stages (MegaScale-MoE):

Tiled pipeline: $\text{dispatch}(i+1) \parallel \text{experts}(i) \parallel \text{combine}(i-1)$



```
recv, recv_work, recv_le, le_work = dispatch(plans[0])
for i in range(n):
    nxt = dispatch(plans[i+1]) if i+1 < n else None    # prefetch next tile's a2a
    recv_work.wait()                                # ... while this tile computes
    out_i = _ffn_forward(recv_sorted, counts, gate_w, up_w, down_w)
    combined, combine_work = all_to_all_no_grad_async(out_i, ...)    # combine in flight
```

Every rank issues exactly N collectives (empty tiles still call the zero-sized a2a) → **lockstep**.

Opt 2 – tiled backward (the hard part)

“ MegaScale-MoE's lesson, re-discovered the hard way:
do not implement the MoE backward with `torch.autograd`. ”

A first version calling `torch.autograd.grad` twice per tile regressed the full step to **0.87×**.

Why: autograd schedules backward nodes **serially** — you cannot make it interleave one tile's reverse-a2a with another tile's compute. **You must own the backward schedule.**

Megatron's independence trick: from the same upstream grad, the FFN backward produces `d_input` and `d_weight` **independently**:

```
per tile: dgrad → launch reverse-dispatch a2a(d_input) || compute wgrad  
           combine_rev(i+1) is prefetched to cover tile i's recompute
```

→ one recompute pass, no double graph traversal, comm hidden behind `wgrad`.

Opt 2 – the Triton megakernel

Production-style fused expert FFN (vLLM / MegaBlocks pattern):

- **Grouped GEMM**: fixed `BLOCK_M=64` token blocks, each block belongs to **one** expert (block-aligned tiling). `triton.autotune` over tile sizes / warps / stages.
- gate & up **fused** into one GEMM `w13 = [E, 2I, H]` → Triton `silu_and_mul` → down GEMM.
- routed-weight epilogue fused into the down-GEMM store.

Explicit backward (no autograd):

dgrad: $dX = dC \cdot W$ (reuse fwd GEMM, swap W strides: $AW^\top \rightarrow AW$)

wgrad: $dW = C^\top A$ (dedicated grouped-GEMM kernel)

```
dx_sorted, a_c, di_c = megakernel.fused_moe_dgrad(recv, counts, gate_w, up_w, down_w, d_out)
wgrads = lambda: megakernel.fused_moe_wgrad(recv, counts, a_c, di_c, d_out) # deferred → overlap
```

Opt 3 – DeepEP backend

Swap our `torch.distributed` a2a for **DeepEP's hand-written CUDA kernels** (Hopper SM90+).

- Saturates NVLink/RDMA at **near-zero SM occupancy** → frees SMs for the GEMM (its real win is the cross-node RDMA path, not raw 2-GPU bandwidth).
- We wrap the **classic intranode Buffer** (NVLink P2P), *not* the V2 `ElasticBuffer` – the latter needs NCCL GIN, which aborts on single-node clusters (e.g. Modal H100): *"NCCL GIN is unavailable"*.

```
recv_x, handle, recv_idx, recv_w, _ = deeppep_backend.dispatch(group, latent, topk_idx, topk_w, E)
expert_out = self._deeppep_experts(recv_x, recv_idx, recv_w) # grouped FFN + gate weights
y = deeppep_backend.combine(group, expert_out, handle) # sum per-rank contributions
```

`dispatch` and `combine` are **linear transposes** ⇒ autograd backward of one **is** the other.
Gate weights ride a non-differentiable transport (router grad flows the normal way).

Always safe: on A100 (sm_80) `deeppep_available()` is False → transparent fallback to torch.

Switching families — reduce *volume*

Overlap hides the a2a. But if $\rho \ll 1$, or you're already hiding everything, the next lever is **fewer bytes on the wire**.

FP8 dispatch

$2\text{ B} \rightarrow \sim 1\text{ B}$ per element.

Comm-bound win.

LatentMoE

route/compute in dim $1 \ll d$.

Cuts comm **and** FLOPs.

These are *architecture/precision* changes to the payload, independent of *how* it's moved — so they stack with each other **and** with all three overlap paths.

Opt 4 – FP8 (E4M3) dispatch

DeepSeek-V3 recipe: **FP8 dispatch, BF16 combine.**

Per-token scale into E4M3's range (max magnitude **448**):

$$s_t = \frac{\max_h |x_{t,h}|}{448}, \quad \hat{x}_t = \text{clip}\left(\frac{x_t}{s_t}, \pm 448\right) \in \text{e4m3}$$


```
q, scale = _per_token_quant_fp8(x)                # [N,H] fp8, [N] fp32
dist.all_to_all_single(recv_q.view(uint8), q.view(uint8), ...) # NCCL has no fp8 coll → bitcast
recv = recv_q.to(bf16) * recv_scale[:, None]        # dequant on receiver, GEMM in BF16
```

- Backward is **straight-through** in BF16 (only the forward *activation* dispatch is low precision).
- No FP8 tensor cores needed → **works on A100**.
- Measured: **~2× fewer dispatch bytes**, relative L2 error **0.023**.

Opt 5 – LatentMoE (NVIDIA Nemotron)

Shared **down-projection** $d \rightarrow 1$ before routing/dispatch; experts & a2a live in 1 ; **up-projection** $1 \rightarrow d$ after combine. Router & shared expert stay in full d .

LatentMoE: route and compute in a small latent dim l (l much less than d)



- It's an **architecture** change (experts parameterized in 1) $\rightarrow$ must be **trained** this way.
- Nemotron reinvests the savings into **more experts / higher top-k** (accuracy-per-FLOP).
- The **only** optimization here that also cuts *compute* $\Rightarrow$ wins even on fast fabrics.

Stacking the volume reducers

hidden=4096, latent=1024, inter=4096, tokens=8192, dispatch payload:

config	wire B / route	vs dense
dense bf16	8192	1.0×
dense + fp8	4100	2.0×
latent (l=1024)	2048	4.0×
latent + fp8	1028	8.0×

The two are multiplicative on bytes. What that buys in wall-clock depends on $\rho \rightarrow$ next.

Results I – PCIe is comm-bound (ρ large)

2× L4, **PCIe**, no NVLink. Single MoE layer (`hidden4096 inter4096 experts8 tokens8192`, dispatch ≈ 134 MB):

path	forward	fwd+bwd
plain	1.00×	1.00×
overlap(shared)	1.10×	1.02×
tiled-3	1.39×	1.25×
tiled-4	1.62×	1.27×
tiled-6	1.54×	1.32×

End-to-end training step (real `Llama`, 8192 tok/layer): **tiled-4** $\rightarrow$ **1.28×** (8636 $\rightarrow$ 11078 tok/s).

$\rightarrow$ When the link is slow, hiding the a2a is a **big** win.

Results II – NVLink is compute-bound (ρ tiny)

2× A100, **NV12** $\approx$ **300 GB/s**. Same code, story **inverts**:

path	fwd (134 MB)	fwd+bwd	fwd (537 MB)
plain	1.00× (21.4 ms)	1.00× (77.4 ms)	1.00× (79.5 ms)
overlap(shared)	1.00×	1.01×	1.01×
tiled-3	0.97×	0.98×	0.98×
tiled-6/8	0.86–0.94×	0.96×	0.91–0.92×

Even a 537 MB payload moves dispatch+combine ($\sim$ 1 GB) in **$\sim$ 3.6 ms** = **$\sim$ 4.5 % of the step**.
Nothing to hide $\rightarrow$ per-tile overhead **regresses** it. **Default** `ep_num_tiles=1` **on NVLink**.

Results III – controlled proof: it's *bandwidth*, not hardware

Same A100 box, same config, forward-only. **Only variable: interconnect bandwidth** (`NCCL_P2P_DISABLE=1` forces the slow host path).

interconnect	plain	best tiled	speedup
NVLink on (~300 GB/s)	21.4 ms	22.0 ms	none (comm $\approx 4.5\%$)
NVLink off (slow path)	39.2 ms	30.5 ms	1.28×
2× L4 PCIe	—	—	1.3–1.6×


Flip the bandwidth, the same kernels flip from no-op to 1.28×

“ Reducing `inter` does **not** reproduce this — it leaves comm fast and only exposes fixed per-tile launch overhead (tiled-6 $\rightarrow 0.43\times$). **Only throttling bandwidth restores the bottleneck.**

”

The picture: speedup tracks bandwidth, not hardware

Same code, three fabrics: overlap speedup tracks $\rho = t_{\text{comm}}/t_{\text{compute}}$



Identical code & config. The overlap win appears **only** where comm is a real fraction of the step — exactly the cross-node / large-EP regime where production MoE lives.

Results IV – H100 + DeepEP

Built from source on **2× H100 (NVLink) via Modal** (`modal run modal_run.py`). All correctness checks pass on real Hopper (DeepEP vs torch, fwd **and** bwd).

model	fwd rel err	grad rel err	torch	deepep	speedup
dense	0.0000	0.0000	2.00 ms	1.64 ms	1.22×
LatentMoE (<code>l=512</code>)	0.0000	0.0000	2.06 ms	1.66 ms	1.24×

DeepEP matches torch to bf16 and is ~1.2× faster **even on one NVLink node** (frees SMs, skips our explicit permute). The big wins are at scale / cross-node.

Compress on H100 (fwd): latent **2.41×**, latent+fp8 2.26× — same story as A100.

Results V – EP scaling: the wall, live (8× H100)

Weak scaling (4 experts/GPU, 8192 tok/GPU fixed; `num_experts = 4·ep`). `slowlink` = `NCCL_P2P_DISABLE=1` (emulates cross-node). Forward, one MoE layer:

ep	NVLink plain	slowlink plain	$\rho \approx \text{comm/comp}$	slowlink best tiled
2	11.2 ms	45.4 ms	3.1	1.17×
4	12.1 ms	64.0 ms	4.3	1.14×
8	12.3 ms	203.7 ms	15.5	1.16×

Per-rank payload is **constant**, yet bandwidth-bound a2a explodes **45→204 ms** (compute flat ~12 ms) → ρ blows up **3→15**: *the communication wall*. Overlap recovers ~1.15× fwd, but at ep=8 **fwd+bwd regresses to 0.88×** — once ρ is double-digit you must **cut comm** (FP8 / LatentMoE / DeepEP), not just hide it.

Beyond speed

Back to the Memory & Compute walls

the comm wall is down — now fixed memory (Memory Wall) + balanced experts (Compute Wall): MAI / GShard / Switch / DeepSeek-V3

Why speed isn't enough

The all-to-all is fast now — but two **router-driven** walls remain, and neither is about latency:

Memory Wall → memory swings

The router decides per-expert counts **every step**. Dropless buffers track them → activation memory fluctuates → allocator **fragmentation & OOM** at scale.

Compute Wall → expert collapse

A trained router with no balancing pressure piles tokens onto a **few** experts → the rest sit idle (wasted GEMMs); capacity (below) then **drops** the overflow.


→ **Opt 6** breaks the Memory Wall (fixed-capacity buffers); **Opt 7** attacks the Compute Wall's load imbalance.

Opt 6 – capacity-capped / static-memory dispatch

Run experts on **fixed buffers**: each `(rank, expert)` processes exactly `c` tokens per round.


Static-memory dispatch: stop sizing buffers to the router

Dropless (naive/tiled): buffer = actual count



heights swing every step -> fragmentation / OOM

Capacity-capped: fixed C per (rank, expert)



fixed shape (fwd+bwd); dropless = extra capped rounds

$$C = \text{factor} \cdot \frac{\text{tokens} \cdot k}{E}, \quad \text{send} \in \mathbb{R}^{E \times C \times d} \rightarrow \text{reshape} [ep, L, C, d] \rightarrow \text{balanced a2a}$$

- **capped** (`dropless=false`): one round, overflow dropped (GShard / Switch).
- **static-memory dropless** (default): rounds `r` process slots `[rC, (r+1)C]` until all tokens are done; round count `MAX`-reduced over EP so ranks stay in lockstep.

Opt 6 – deep notes (all autograd, fixed bwd too)

The fixed `[E, C, d]` buffer makes the EP exchange a **perfectly balanced** all-to-all of equal `L·C` blocks — and the whole path is plain autograd, so the **backward also runs on fixed-capacity buffers** (no custom `Function`).


```
target = expert_idx[keep] * C + slot[keep]           # fixed-capacity addressing
send   = zeros(E*C, d).index_add(0, target, routed[keep]) # scatter (differentiable)
recv   = all_to_all(send, [L*C]*ep, [L*C]*ep, group)  # balanced, fixed-shape
out    = experts(recv.reshape(L, ep*C, d))           # uniform counts -> grouped GEMM
contrib = all_to_all(out, ...)[target] * weights[keep] # gather + gate (differentiable)
```

Droptless + enough rounds = bit-exact with naive (`out_diff = 0`, `grad_diff ≤ 3e-7`, incl. LatentMoE; `tests/test_capacity_moe.py`).


Opt 7 – load balancing the router

Gate **weights** stay unbiased (forward never skewed). Two composable controllers:

Load balancing: keep every expert busy (per-expert token load)



Loss-free bias + aux loss



aux loss: $\alpha * \sum_e f_e * P_e$ (grad to gate) | bias: $b_e += \text{rate} * \text{sign}(\text{mean_load} - \text{load}_e)$ (no grad)

Aux loss (Switch / DeepSeek) – $L_{\text{aux}} =$

$$\alpha \sum_e f_e P_e$$

differentiable → grad to the gate;

`collect_aux_loss` adds it to the loss.

Loss-free bias (DeepSeek-V3) – $b_e +=$

$$\eta \text{sign}(\bar{L} - L_e)$$

biases **selection only**: no grad, no tug-of-war with the LM loss.

Opt 6 + 7 – the ablation (`tests/ablation_moe.py`)

Capacity (C fixed at 128, $k=2$): imbalance makes the naive buffer grow while C stays put; dropless stays **bit-exact**.

skew	naive buf	capped drop	dropless rounds
0.0	144	2.6%	2
1.0	320	23.9%	3
2.0	459	47.8%	4

`factor` 0.5→2.0 trades buffer for drops: 50.8% → 6.2%.

Load balancing (400 steps, 8 clusters): bias revives the dead expert; aux stacks.

config	MaxVio	dead
none	1.51	1
aux	1.49	0
bias	1.05	0
aux+bias	1.05	0

MaxVio = busiest / mean load (1.0 = perfect).

Opt 6 on GPU – static memory, measured (4× H100)

`bench_capacity_moe.sh`: one MoE layer fwd+bwd, `experts=16 tok/GPU=8192 hidden=4096`, `C=1024`.

Peak CUDA memory (MAX over ranks) as routing imbalance grows:

skew	max/mean	uncapped peak / ms	capped-drop peak / ms	dropless peak / ms
0.0	1.03	4790 MB / 37	4698 MB / 36	5435 MB / 72
1.0	4.15	9508 MB / 83	4698 MB / 36	7046 MB / 176
2.0	7.10	13025 MB / 118	4698 MB / 36	8657 MB / 288

Uncapped peak **swings 4.8→13 GB (2.7×)** with the router – the memory wall, live.

`capped-drop` is **dead flat** (4698 MB / 36 ms) but drops up to 60% at `factor=1`; `dropless` keeps every token and still caps far below uncapped, paying rounds in latency. At balanced load (skew 0) capacity is a slight *loss* – it's an at-scale **safety/memory** lever, not a free speedup.

The big lesson

Overlap (3 paths)

Bandwidth-bound. Pays off iff $\rho = t_{\text{comm}}/t_{\text{compute}}$ is non-trivial:

- cross-node IB/RoCE

- large EP degree

- single-node NVLink


LatentMoE

Compute-bound win too. Cuts FLOPs, so it pays off **even on NVLink** (2.14×).
The one lever that doesn't need a comm bottleneck.

FP8 dispatch

Comm-bound like overlap (+7% off-NVLink, slight regression on NVLink).

Production MoE lives in the **cross-node, large-EP** regime → that's the `NCCL_P2P_DISABLE=1` row, where everything here compounds.

Decision guide

Your situation	Turn on
Single-node NVLink, few tokens/layer	<code>ep_num_tiles=1</code> , <code>ep_overlap</code> ~free, skip FP8
Single-node NVLink, want speed	LatentMoE (cuts compute)
Cross-node / IB / large EP, $\geq 8k$ tok/layer	tiled-N (+ overlap)
Hopper + want SMs back	<code>ep_backend="deepep"</code>
Slow fabric, comm-bound	FP8 dispatch + LatentMoE (8× bytes)
Memory swings / OOM from imbalance	<code>ep_capacity_factor>0</code> (dropless = bit-exact, fixed memory)
Router collapsing onto few experts	<code>router_aux_loss_coef</code> + <code>router_bias_update_rate</code>
Always	leave knobs on — all auto-fallback & bit-exact

```
"ep_overlap": true, "ep_num_tiles": 4, "ep_fp8_dispatch": false,  
"moe_latent_dim": 1024, "ep_backend": "deepep",  
"ep_capacity_factor": 1.25, "router_aux_loss_coef": 1e-2, "router_bias_update_rate": 1e-3
```

Recap — the optimization ladder

1. **Naive** a2a dispatch/combine — differentiable, reverse-a2a backward.
2. **Shared-expert overlap** — side stream, composes with every path.
3. **Tiled pipeline** — fwd $\parallel$; backward needs **explicit kernels** (autograd-twice = 0.87 $\times$).
4. **Megakernel** — grouped GEMM; dgrad by stride-swap, dedicated wgrad.
5. **DeepEP** — Hopper kernels, near-zero SM, classic **Buffer**, auto-fallback.
6. **FP8 dispatch** — E4M3 per-token scale, straight-through BF16 backward, 2 $\times$ bytes.
7. **LatentMoE** — route/compute in **1**, cuts comm **and** FLOPs, stacks to 8 $\times$ with FP8.
8. **Static-memory dispatch** — fixed-capacity buffers; dropleass rounds = bit-exact, constant memory.
9. **Load balancing** — aux loss + loss-free bias keep every expert busy (router quality).

Throughline: measure ρ . Hide comm where it's expensive; cut bytes/FLOPs where it isn't; then make it **trainable** — fixed memory, balanced experts.

Appendix

file map • how to run

File map & how to run

File	Role
<code>expert_parallel.py</code>	<code>MoELayer</code> : router, sharded experts, path selection, FP8 + LatentMoE
<code>ep_communications.py</code>	differentiable / async / FP8 all-to-all primitives
<code>tiled_moe.py</code>	<code>_TiledMoEFunction</code> : tiled pipeline, fwd and bwd overlap
<code>capacity_moe.py</code>	fixed-capacity static-memory dispatch (capped + dropleass rounds)
<code>megakernel.py</code>	Triton fused FFN + explicit <code>dgrad</code> / <code>wgrad</code> grouped GEMM
<code>deepep_backend.py</code>	DeepEP dispatch/combine (Hopper, auto-fallback)

```
# correctness (2 ranks, gloo/CPU by default; EP_TEST_BACKEND=ncc1 for GPU)
torchrun --nproc_per_node 2 tests/test_expert_parallel.py
torchrun --nproc_per_node 2 tests/test_capacity_moe.py
python tests/test_load_balance.py
torchrun --nproc_per_node 2 tests/test_fp8_dispatch.py

# benchmarks (2 GPUs)
torchrun --nproc_per_node 2 tests/bench_ep_overlap.py --tokens 8192 --tiles 3 4 6 --backward
torchrun --nproc_per_node 2 tests/bench_ep_compress.py --hidden 4096 --latent 1024 --tokens 8192

# DeepEP / Hopper on cloud GPUs (builds DeepEP from source)
modal run modal_run.py
```