

Tensor Parallelism, from scratch

What ships in `picotron`, and three improvements I layered on top

Megatron column/row split → **async overlap** · **sequence parallelism** · **vocab-parallel cross-entropy**

all bit-exact with the baseline, gradient-validated and benchmarked on 2× A100

What we are parallelizing


A transformer block is two GEMM sandwiches: **attention** (`q,k,v → attn → out`) and **MLP** (`up,gate → act → down`). Tensor parallelism splits those matmuls *inside* the layer across the TP group.

- splits **within** a layer $\Rightarrow$ a collective **every** layer $\Rightarrow$ lives on NVLink, TP degree kept small (≤ 8)
- (contrast: pipeline splits **across** layers, one neighbor hop per stage boundary)

Megatron-LM, Shoeybi et al. 2019 – [arXiv:1909.08053](https://arxiv.org/abs/1909.08053)

The column → row split

Tensor parallelism: column then row split inside the attention / MLP block



one all-reduce in the forward (g) and one in the backward (f) per block; the head / intermediate dim stays sharded throughout

- **Column-parallel** (q/k/v, up/gate): shard by output columns. Input replicated, output sharded.
- **Row-parallel** (out, down): shard by input rows. Input already sharded, output is a partial sum.

Chaining col→row ⇒ **one all-reduce in fwd, one in bwd** per block; heads/intermediate stay sharded.

The **f** and **g** operators

Two conjugate autograd functions — that's the whole communication contract.

```
class CopyToModelParallelRegion(Function):    # f: column-parallel input
    def forward(ctx, x): return x              # copy
    def backward(ctx, g): all_reduce(g); return g    # all-reduce

class ReduceFromModelParallelRegion(Function): # g: row-parallel output
    def forward(ctx, x): all_reduce(x); return x    # all-reduce
    def backward(ctx, g): return g                 # copy
```

The embedding is sharded over the **vocabulary** (`VocabParallelEmbedding`): mask off-shard tokens,
look up, all-reduce. This is the correct Megatron-v1 TP that already ships in picotron.

Three improvements, one ladder

Each is **opt-in** (defaults unchanged) and **bit-exact** with the baseline.

rung	does	buys
+async	overlap col-parallel input-grad all-reduce with the weight-grad GEMM	hide comm
+seqpar	shard norm / residual along the sequence	tp× less activation memory, <i>same bytes</i>
+vocab_ce	keep logits vocab-sharded; loss sends [b, s] scalars	no [b, s, V] logits, $O(V) \rightarrow O(1)$ comm

```
apply_tensor_parallel(model, async_tp=True, sequence_parallel=True, vocab_parallel_ce=True)
```

Rung 1 — async communication overlap

The column-parallel backward owes the previous layer $\text{grad_x} = \text{grad_y} \cdot W$ (all-reduced) **and** the

local weight grad. Launch the grad_x all-reduce **first**, compute grad_W while it flies:


```
grad_input = grad_output @ weight
handle = dist.all_reduce(grad_input, async_op=True)  # launch, don't wait
grad_weight = grad_output.t() @ input_             # overlaps the collective
handle.wait()
```

- The code existed but was never wired in — `async_tp=True` turns it on.
- Needs `CUDA_DEVICE_MAX_CONNECTIONS=1` so the collective and GEMM actually run concurrently.

Rung 2 – sequence parallelism

Between blocks (norm + residual) activations are **replicated** → that memory ignores TP.
Shard them
along the **sequence** instead:

Sequence parallelism: shard the norm / residual regions along the sequence



all-reduce = reduce-scatter + all-gather => same bytes on the wire, but the stream is tp x smaller
the fused col-parallel linear checkpoints only the s/tp input and re-gathers in backward; that is what frees the memory

The identity that makes it free on bytes: **all-reduce = reduce-scatter + all-gather.**

Rung 2 — the subtlety that decides the win

A naive `all_gather` + `F.linear` **loses** the memory win: autograd saves the *gathered* full-sequence

input for the weight-grad, so the biggest activation is still full-size on every rank.

Fix (Megatron): a fused gather+linear that checkpoints only the **sharded** input and re-gathers in backward.

```
class _ColumnParallelLinearWithSequenceParallel(Function):
    def forward(ctx, x_sharded, w, b):
        ctx.save_for_backward(x_sharded, w)          # NOT the gathered tensor
        return all_gather_seq(x_sharded) @ w.t() + b
    def backward(ctx, g):
        x_full = all_gather_seq(x_sharded)            # recompute, don't store
        ...                                           # reduce-scatter grad_x; grad_w from x_full
```

“ This one change turned the measured SP saving from **51 MB → 3.3 GB**. ”


cost: each rank's RMSNorm sees a different seq shard $\Rightarrow$ all-reduce its weight grad over TP (a "sequence-parallel" param).

Rung 3 – vocab-parallel cross-entropy

The output projection is already column-parallel over the vocab. The stock path all-gathers the full

$[b, s, V]$ logits and runs a dense softmax on every rank – the biggest activation in the step.

Vocab-parallel cross-entropy: never gather the $[b, s, V]$ logits



The full logits never exist on any rank; for a 32k-128k vocab this is the single biggest activation saved.

Keep logits sharded; exchange only per-token **max**, **$\Sigma \exp$** , **target logit**. $O(V) \rightarrow O(1)$ comm, full logits never materialized. Backward is the saved-softmax **softmax - onehot(target)**.

Bit-exact — every rung

```
# full Llama: plain TP and TP+sequence-parallel vs a single-GPU reference
torchrun --nproc_per_node 2 tests/test_tp_sequence_parallel.py
# vocab-parallel CE vs dense F.cross_entropy on gathered logits
torchrun --nproc_per_node 2 tests/test_tp_vocab_ce.py
```

check	result
plain TP vs single-GPU reference	loss_diff = 0, grad_diff $\approx 2\text{e-}8$
TP + sequence parallel vs reference	loss_diff = 0, grad_diff $\approx 2\text{e-}8$
vocab-parallel CE vs dense CE	grad_diff $\approx 4\text{e-}9$

The sequence-parallel test also all-reduces the norm weight grads over TP — exactly what real training must do.

The ablation (2× A100, TP=2, bf16)

24-layer · hidden 2048 · seq 2048 · mbs 4 · SDPA (`tests/bench_tp.py`):

config	ms/step	tok/s	peak MB	mem vs base	speedup
baseline	327.7	49992	17865	1.00×	1.00×
+async	320.5	51126	17865	1.00×	1.02×
+seqpar	363.2	45107	14576	1.23×	0.90×
+vocab_ce	321.7	50936	15986	1.12×	1.02×
+seqpar+vocab_ce	356.4	45966	12697	1.41×	0.92×

2nd shape (L16/h3072/seq4096/vocab49k): seqpar 1.17× · vocab_ce 1.14× · both **1.37×** — same story.

The honest lesson

Sequence parallelism

A **memory** lever, not a speed one. $1.23\times$ less memory by sharding the residual stream; the extra reduce-scatter / all-gather *calls* cost $\sim 10\%$ on a single NVLink node at TP=2. Pays off more at higher TP / across nodes (and unlocks not recomputing those regions).

They stack: $1.41\times$ less memory → a longer sequence or bigger micro-batch on the same card.

Vocab-parallel CE

Nearly free — $1.12\text{--}1.14\times$ memory at neutral speed. Kills the `[b,s,v]` gather + dense softmax for a few `[b,s]` all-reduces; helps more as the vocab grows.

Async overlap

Small win here ($1.02\times$); grows with the weight-grad GEMM vs all-reduce latency.

Decision guide

Your situation	Turn on
Always (it's nearly free, big vocab)	<code>vocab_parallel_ce=True</code>
Tight on activation memory / want longer seq	<code>sequence_parallel=True</code>
Large hidden, comm-latency visible, single connection	<code>async_tp=True</code> + <code>CUDA_DEVICE_MAX_CONNECTIONS=1</code>
Memory-bound at scale (high TP, cross-node)	seqpar + vocab_ce together

```
apply_tensor_parallel(model, async_tp=False, sequence_parallel=True, vocab_parallel_ce=True)
```

Appendix

file map • how to run

File map & how to run

File	Role
<code>tensor_parallel.py</code>	<code>ColumnParallelLinear</code> , <code>RowParallelLinear</code> , <code>VocabParallelEmbedding</code> , <code>apply_tensor_parallel</code>
<code>tp_communications.py</code>	<code>f / g</code> collectives, <code>async-overlap linear</code> , <code>vocab_parallel_cross_entropy</code>
<code>../sequence_parallel/</code>	seq-parallel collectives + fused gather-linear (shares the TP group)
<code>tests/test_tensor_parallel.py</code>	column/row vs reference <code>nn.Linear</code>
<code>tests/test_tp_sequence_parallel.py</code>	plain TP & TP+SP vs single-GPU reference (bit-exact)
<code>tests/test_tp_vocab_ce.py</code>	vocab-parallel CE vs dense cross-entropy
<code>tests/bench_tp.py</code>	the ablation: step time / tokens/s / peak memory per rung

```
CUDA_DEVICE_MAX_CONNECTIONS=1 torchrun --nproc_per_node 2 tests/test_tp_sequence_parallel.py
CUDA_DEVICE_MAX_CONNECTIONS=1 torchrun --nproc_per_node 2 tests/bench_tp.py \
  --hidden 2048 --inter 5632 --layers 24 --vocab 32768 --seq 2048 --mbs 4
```